




**CSCS**

Centro Svizzero di Calcolo Scientifico  
Swiss National Supercomputing Centre

**ETH** zürich



# Multi-Node Distributed Training with tf.keras

Summer School 2020

Henrique Mendonça, CSCS

July 24, 2020

# Table of Contents

- Motivation
- Distributed training with tf.keras 2.2
- Multi-worker/node distribution
- Multi-worker distribution with SLURM
- Batch Norm Synchronisation
- Scaling learning rate and momentum
- Assignment



# Motivation

- Why should we use multiple GPU's to train a DL model?



# Motivation

- Because it's faster
  - Allow quick experimentation of new ideas and models




- Allows training models larger than memory
  - Out-of-Core learning: Not covered in this course
- May allow better accuracy, especially in larger models. Why?


# Motivation

Visualizing the Loss Landscape of Neural Nets

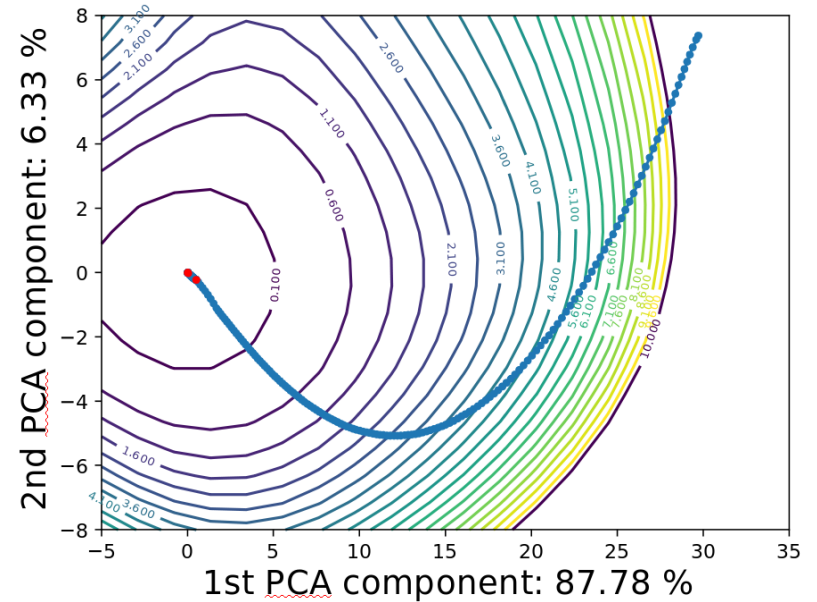
<https://arxiv.org/pdf/1712.09913.pdf>



Batch size 128

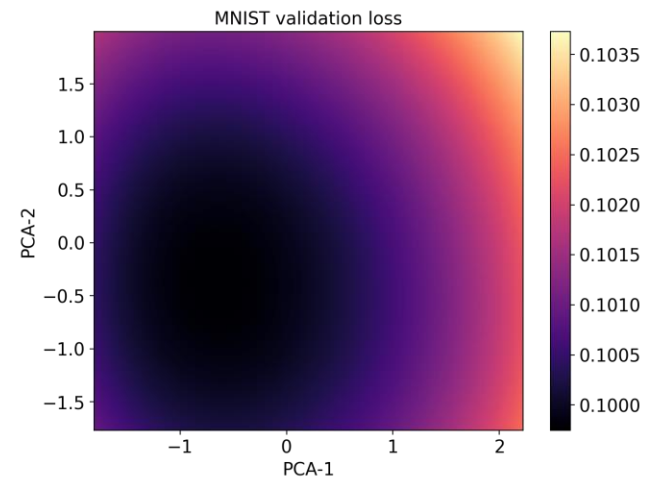



Batch size 8192


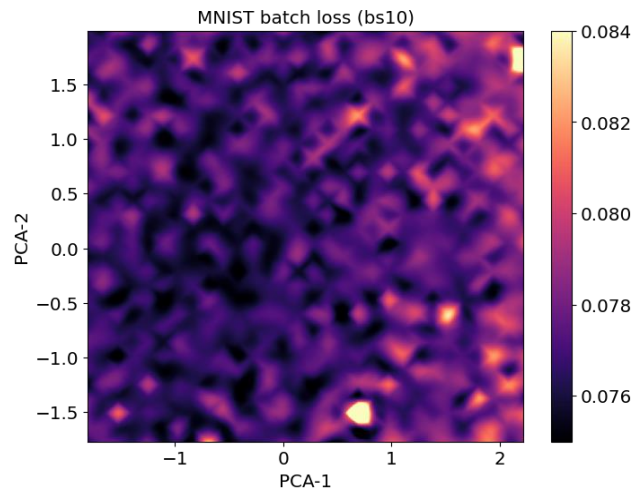


# Motivation

- Validation:



- Training:



# Distributed training with TF Keras 2.2

- TensorFlow introduces [tf.distribute.Strategy](#)
  - Wraps model in a distributed scope
- Multi-GPU -> [tf.distribute.MirroredStrategy](#)
  - Single node/worker

```
strategy = tf.distribute.MirroredStrategy()  
with strategy.scope():  
    model = build_and_compile_model()  
model.fit(dataset, epochs, steps_per_epoch)
```

- [NCCL](#) AllReduce by default
- Automatic data sharding across GPU's

# Multi-worker Distribution

- Creates some additional complexity
  - external network communication
  - separated OS
  - separated processes
    - Facilitated by `ipyparallel` magic on JupyterLab





# Multi-worker Distribution

- [tf.distribute.experimental.MultiWorkerMirroredStrategy](#)
- Communication:
  - NCCL AllReduce for all-reduce (if available)
  - Ring algorithm for all-gather
  - Includes fault tolerance when using [ModelCheckpoint](#)\*
- Cluster Resolver:
  - defaults to TFConfig

```
os.environ['TF_CONFIG'] = '{  
    "cluster": {"worker": ["nid01111:8888", "nid02222:8888"]},  
    "task": {"type": "worker", "index": "0"}  
}'
```

\* TF 2.3 uses the [BackupAndRestore](#) callback instead

# Multi-worker distribution with SLURM

- TensorFlow 2.2+

```
tf.distribute.cluster_resolver.SlurmClusterResolver(  
    port_base=8888, auto_set_gpu=True, rpc_layer='grpc',  
    jobs=None, gpus_per_node=None, gpus_per_task=None,  
    tasks_per_node=None  
)
```

- All parameters are automatically queried from SLURM



# Multi-worker distribution with SLURM

```
%%px
strategy = tfd.experimental.MultiWorkerMirroredStrategy(
    cluster_resolver=tfd.cluster_resolver.SlurmClusterResolver(),
    communication=tfd.experimental.CollectiveCommunication.NCCL,
)
with strategy.scope():
    model = build_and_compile_model()
model.fit(dataset, epochs, steps_per_epoch)
```

# Done!



# Multi-worker distribution with SLURM

- Practice
  - Run MNIST training and inference on 2 GPU's

# Batch Norm Synchronisation

- How does Batch Normalisation work?
- What happens if it's used with small batch sizes?

# Batch Norm Synchronisation

- BN: Exponential Moving Average per channel

$$\mu_B = \frac{1}{m} \sum_{i=1}^m x_i$$

$$\sigma_B^2 = \frac{1}{m} \sum_{i=1}^m (x_i - \mu_B)^2$$

$$\hat{x}_i^{(k)} = \frac{x_i^{(k)} - \mu_B^{(k)}}{\sqrt{\sigma_B^{(k)2} + \epsilon}}$$

$$y_i^{(k)} = \gamma^{(k)} \hat{x}_i^{(k)} + \beta^{(k)}$$

- [tf.keras.layers.experimental.SyncBatchNormalization](#)
  - AllReduce across BN layers during forward pass
  - Synchronizes all statistics



**CSCS**

Centro Svizzero di Calcolo Scientifico  
Swiss National Supercomputing Centre

**ETH** zürich



## Assignment

---

### Kaggle: SIIM-ISIC Melanoma Classification

Identify melanoma in lesion images

[kaggle.com/c/siim-isic-melanoma-classification](https://kaggle.com/c/siim-isic-melanoma-classification)

# Dataset - SIIM-ISIC Melanoma Classification

	image_name	patient_id	sex	age_approx	anatom	diagnosis	target
0	ISIC_2637011	IP_7279968	male	45	head/neck	unknown	0
1	ISIC_0015719	IP_3075186	female	45	upper extremity	unknown	0
2	ISIC_0052212	IP_2842074	female	50	lower extremity	nevus	0
3	ISIC_0068279	IP_6890425	female	45	head/neck	melanoma	1
4	ISIC_0074268	IP_8723313	female	55	upper extremity	unknown	0
...	...	...	...	...	...	...	...
33121	ISIC_9999134	IP_6526534	male	50	torso	unknown	0
33122	ISIC_9999320	IP_3650745	male	65	torso	melanoma	1
33123	ISIC_9999515	IP_2026598	male	20	lower extremity	unknown	0
33124	ISIC_9999666	IP_7702038	male	50	lower extremity	unknown	0
33125	ISIC_9999806	IP_0046310	male	45	torso	nevus	0
33126 rows							




# SIIM-ISIC Melanoma Classification

- Practice
  - Train on your 2 GPU's
  - Experiment with hyper parameters / another model
  - Submit to competition

# SIIM-ISIC Melanoma Classification

- Homework
  - Try Kaggle and Colab's free TPU's
    - [tf.distribute.experimental.TPUStrategy](#)
  - Read competition discussion
    - Try-out your own ideas
    - Ensemble
    - Win a Kaggle competition medal



# Thank you for your attention.