




CSCS

Centro Svizzero di Calcolo Scientifico
Swiss National Supercomputing Centre

ETH zürich



Multi-Node Distributed Training with tf.keras

Summer School 2020

Henrique Mendonça, CSCS

July 24, 2020

Table of Contents

- Motivation
- Distributed training with tf.keras 2.2
- Multi-worker/node distribution
- Multi-worker distribution with SLURM
- Batch Norm Synchronisation
- Scaling learning rate and momentum
- Assignment



Motivation

- Why should we use multiple GPU's to train a DL model?



Motivation

- Because it's faster
 - Allow quick experimentation of new ideas and models




- Allows training models larger than memory
 - Out-of-Core learning: Not covered in this course
- May allow better accuracy, especially in larger models. Why?


Motivation

Visualizing the Loss Landscape of Neural Nets

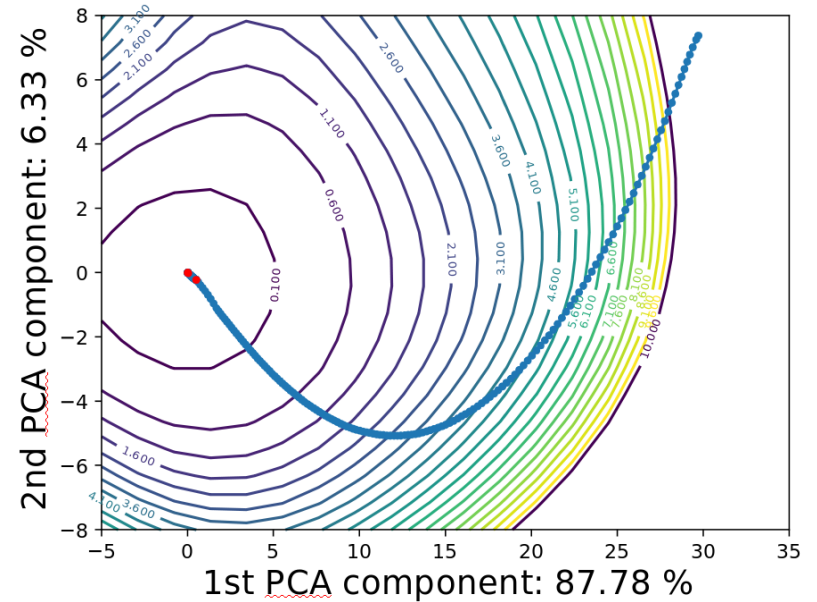
<https://arxiv.org/pdf/1712.09913.pdf>



Batch size 128

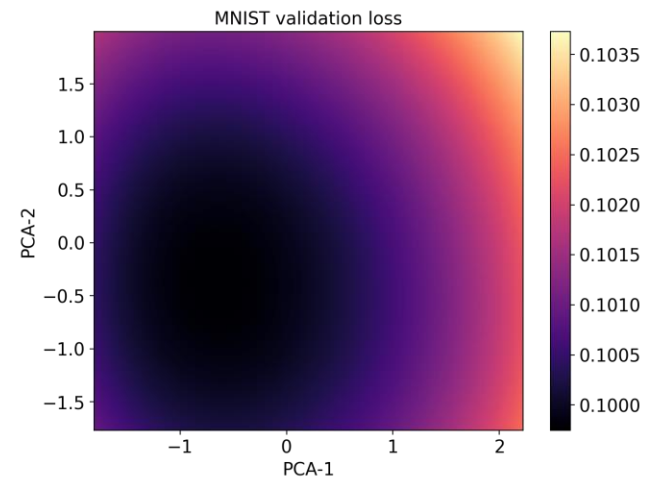



Batch size 8192


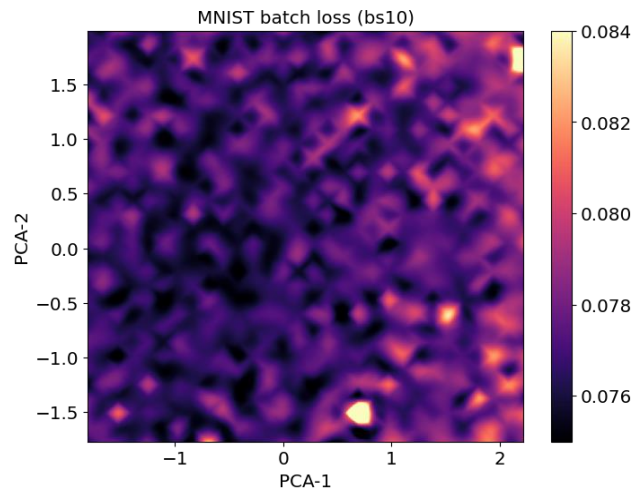


Motivation

- Validation:



- Training:



Distributed training with TF Keras 2.2

- TensorFlow introduces [tf.distribute.Strategy](#)
 - Wraps model in a distributed scope
- Multi-GPU -> [tf.distribute.MirroredStrategy](#)
 - Single node/worker

```
strategy = tf.distribute.MirroredStrategy()  
with strategy.scope():  
    model = build_and_compile_model()  
model.fit(dataset, epochs, steps_per_epoch)
```

- [NCCL](#) AllReduce by default
- Automatic data sharding across GPU's

Multi-worker Distribution

- Creates some additional complexity
 - external network communication
 - separated OS
 - separated processes
 - Facilitated by `ipyparallel` magic on JupyterLab



Multi-worker Distribution

- [tf.distribute.experimental.MultiWorkerMirroredStrategy](#)
- Communication:
 - NCCL AllReduce for all-reduce (if available)
 - Ring algorithm for all-gather
 - Includes fault tolerance when using [ModelCheckpoint](#)*
- Cluster Resolver:
 - defaults to TFConfig

```
os.environ['TF_CONFIG'] = '{  
    "cluster": {"worker": ["nid01111:8888", "nid02222:8888"]},  
    "task": {"type": "worker", "index": "0"}  
}'
```

* TF 2.3 uses the [BackupAndRestore](#) callback instead

Multi-worker distribution with SLURM

- TensorFlow 2.2+

```
tf.distribute.cluster_resolver.SlurmClusterResolver(  
    port_base=8888, auto_set_gpu=True, rpc_layer='grpc',  
    jobs=None, gpus_per_node=None, gpus_per_task=None,  
    tasks_per_node=None  
)
```

- All parameters are automatically queried from SLURM



Multi-worker distribution with SLURM

```
%%px
strategy = tfd.experimental.MultiWorkerMirroredStrategy(
    cluster_resolver=tfd.cluster_resolver.SlurmClusterResolver(),
    communication=tfd.experimental.CollectiveCommunication.NCCL,
)
with strategy.scope():
    model = build_and_compile_model()
model.fit(dataset, epochs, steps_per_epoch)
```

Done!



Multi-worker distribution with SLURM

- Practise
 - Run MNIST training and inference on 2 GPU's

Batch Norm Synchronisation

- How does Batch Normalisation work?
- What happens if it's used with small batch sizes?

Batch Norm Synchronisation

- BN: Exponential Moving Average per channel

$$\mu_B = \frac{1}{m} \sum_{i=1}^m x_i$$

$$\sigma_B^2 = \frac{1}{m} \sum_{i=1}^m (x_i - \mu_B)^2$$

$$\hat{x}_i^{(k)} = \frac{x_i^{(k)} - \mu_B^{(k)}}{\sqrt{\sigma_B^{(k)^2} + \epsilon}}$$

$$y_i^{(k)} = \gamma^{(k)} \hat{x}_i^{(k)} + \beta^{(k)}$$

- [tf.keras.layers.experimental.SyncBatchNormalization](#)
 - AllReduce across BN layers during forward pass
 - Synchronizes all statistics before autodiff

Batch Norm Synchronisation

- Practise
 - try [tf.keras.layers.experimental.SyncBatchNormalization](#)
 - Should we see any accuracy improvement in MNIST?

Scaling learning rate and momentum

- Linear scaling LR
 - One weird trick for parallelizing convolutional neural networks
<https://arxiv.org/pdf/1404.5997.pdf>
 - Accurate, Large Minibatch SGD: Training ImageNet in 1 Hour
<https://arxiv.org/pdf/1706.02677.pdf>
- Square root scaling LR
 - Large Batch Optimization for Deep Learning: Training BERT in 76 minutes
<https://arxiv.org/pdf/1904.00962.pdf> (LAMB optimizer extends LARS)
- Momentum:
 - Smooths out the error surface
 - Very large batches might not require momentum
 - e.g. RMSprop instead of Adam



CSCS

Centro Svizzero di Calcolo Scientifico
Swiss National Supercomputing Centre

ETH zürich



Assignment

Kaggle: SIIM-ISIC Melanoma Classification

Identify melanoma in lesion images

(Final submission deadline: **August 17, 2020**)

kaggle.com/c/siim-isic-melanoma-classification


Dataset - SIIM-ISIC Melanoma Classification

	image_name	patient_id	sex	age_approx	anatom	diagnosis	target
0	ISIC_2637011	IP_7279968	male	45	head/neck	unknown	0
1	ISIC_0015719	IP_3075186	female	45	upper extremity	unknown	0
2	ISIC_0052212	IP_2842074	female	50	lower extremity	nevus	0
3	ISIC_0068279	IP_6890425	female	45	head/neck	melanoma	1
4	ISIC_0074268	IP_8723313	female	55	upper extremity	unknown	0
...
33121	ISIC_9999134	IP_6526534	male	50	torso	unknown	0
33122	ISIC_9999320	IP_3650745	male	65	torso	melanoma	1
33123	ISIC_9999515	IP_2026598	male	20	lower extremity	unknown	0
33124	ISIC_9999666	IP_7702038	male	50	lower extremity	unknown	0
33125	ISIC_9999806	IP_0046310	male	45	torso	nevus	0
33126 rows							


SIIM-ISIC Melanoma Classification

- Transfer Learning made Deep Learning accessible

ImageNet-2012: 1,281,167 samples



1st convolutional layer of
SeResNeXT101.32x4d (ImageNet)




SIIM-ISIC Melanoma Classification

- Practise
 - Train on your 2 GPU's
 - Experiment with hyper parameters
 - Submit to competition

SIIM-ISIC Melanoma Classification

- Homework
 - Try Kaggle and Colab's free TPU's
 - [tf.distribute.experimental.TPUStrategy](#)
 - Read competition discussion
 - Understand the problem and metric
 - Try-out your own ideas
 - Ensemble model predictions
 - Earn a Kaggle competition medal 🏆



Thank you for your attention.