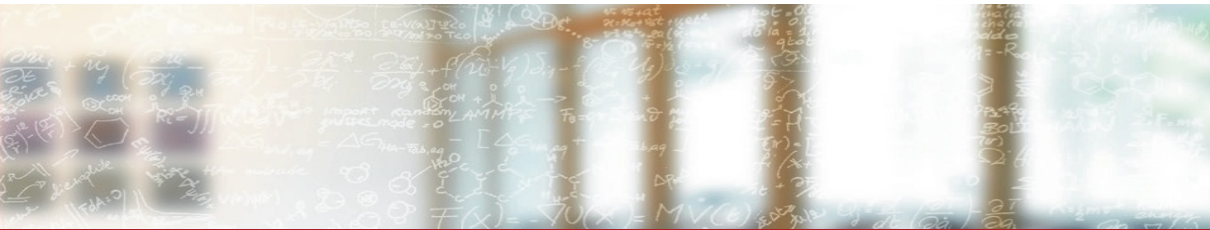




**CSCS**

Centro Svizzero di Calcolo Scientifico  
Swiss National Supercomputing Centre

**ETH** zürich



## Introduction to the Piz Daint environment

CSCS-USI Summer School 2020

*Vasileios Karakasis, CSCS*

July 13, 2020

# Overview

- Accessing CSCS
- Compiling my code
- Running my code
- Editing my code
- Transferring files from/to CSCS
- Repository of the course

# Piz Daint

## Computing nodes

Piz Daint is a hybrid cluster of Cray XC40/XC50 nodes

- Hybrid nodes (XC50)
  - 5320 total
  - Intel® Xeon® E5-2690 v3 @ 2.60GHz (12 cores, 64GB RAM, Haswell)
  - NVIDIA® Tesla® P100 16GB (Pascal)
- Multicore nodes
  - 1813 nodes
  - Two Intel® Xeon® E5-2695 v4 @ 2.10GHz (2 x 18 cores, 64/128 GB RAM, Broadwell)
- Login nodes
  - 5 total
  - Intel® Xeon® CPU E5-2650 v3 @ 2.30GHz (10 cores, 256 GB RAM, Haswell)
- Aries routing and communications ASIC, and Dragonfly network topology

# Piz Daint

## Filesystems

- `/scratch`: High performance Lustre filesystem accessible from the computing nodes
  - Environment variable `$SCRATCH` points to it
  - Total capacity: 6.2 PB
  - Must be used for heavy I/O
- `/users`: GPFS filesystem for the users' homes
- `/project`, `/store`: Long-term storage for computational projects

More on [https://user.cscs.ch/storage/file\\_systems/](https://user.cscs.ch/storage/file_systems/)

# Accessing Piz Daint

- Accessible through SSH
- Piz Daint is not directly accessible from the outside world:
  - `ela` → `daint10x` → `nidxxxxxx`

Two-steps process:

1. Login to the frontend, forwarding X11 (will be needed the second day)
2. Move to the login nodes of Piz Daint

```
# Login to the frontend first
ssh -Y classXXX@ela.cscs.ch
ssh -Y daint
```

# Programming Environments

## Cray Linux Programming Environment

- 4 compilers available: CCE, GNU, INTEL, PGI
- 4 predefined Programming Environments:
  - `PrgEnv-cray` (default), `PrgEnv-gnu`, `PrgEnv-intel`, `PrgEnv-pgi`
  - `echo $PE_ENV` to get the current `PrgEnv`
- 3 wrappers available: `ftn` (Fortran), `cc` (C), `CC` (C++)
  - Required for compiling MPI programs
  - They set appropriate optimisation flags for the target architecture (CPU or GPU)
  - They provide a sort of portability across the programming environments

# Managing programming environments

Daint uses *Environment Modules* (TMod) for managing the programming environments and the software packages:

- Dynamic modification of a user's environment via *modulefiles*.
- All programming environments and software on Daint is available through modules.
- The compiler wrappers will detect the loaded programming environment and automatically set the correct flags and libraries.

# Managing programming environments

## Listing modules

```
$ module list
Currently Loaded Modulefiles:
 1) modules/3.2.11.3
 2) cce/9.0.2
 3) craype-network-aries
 4) craype/2.6.1
 5) cray-libsci/19.06.1
 6) udreg/2.3.2-7.0.1.1_3.9__g8175d3d.ari
 7) ugni/6.0.14.0-7.0.1.1_7.10__ge78e5b0.ari
 8) pmi/5.0.14
 9) dmapp/7.1.1-7.0.1.1_4.8__g38cf134.ari
10) gni-headers/5.0.12.0-7.0.1.1_6.7__g3b1768f.ari
11) xpmem/2.2.19-7.0.1.1_3.7__gdcf436c.ari
12) job/2.2.4-7.0.1.1_3.8__g36b56f4.ari
13) dvs/2.12_2.2.151-7.0.1.1_5.6__g7eb5e703
14) alps/6.6.56-7.0.1.1_4.10__g2e60a7e4.ari
15) rca/2.2.20-7.0.1.1_4.9__g8e3fb5b.ari
16) atp/2.1.3
17) perftools-base/7.1.1
18) PrgEnv-cray/6.0.5
19) cray-mpich/7.7.10
20) slurm/20.02.2-1
21) craype-haswell
22) xalt/2.7.24
```



# Managing programming environments

## Switching programming environments

```
$ module switch PrgEnv-cray/6.0.5 PrgEnv-pgi
```

```
$ module list
```

```
Currently Loaded Modulefiles:
```

- |  |  |
|--|--|
| 1) modules/3.2.11.3                          | 12) dmapp/7.1.1-7.0.1.1_4.8__g38cf134.ari          |
| 2) pgi/19.7.0                                | 13) gni-headers/5.0.12.0-7.0.1.1_6.7__g3b1768f.ari |
| 3) craype-haswell                            | 14) xpmem/2.2.19-7.0.1.1_3.7__gdcf436c.ari         |
| 4) craype-network-aries                      | 15) job/2.2.4-7.0.1.1_3.8__g36b56f4.ari            |
| 5) craype/2.6.1                              | 16) dvs/2.12_2.2.151-7.0.1.1_5.6__g7eb5e703        |
| 6) cray-mpich/7.7.10                         | 17) alps/6.6.56-7.0.1.1_4.10__g2e60a7e4.ari        |
| 7) slurm/20.02.2-1                           | 18) rca/2.2.20-7.0.1.1_4.9__g8e3fb5b.ari           |
| 8) xalt/2.7.24                               | 19) atp/2.1.3                                      |
| 9) udreg/2.3.2-7.0.1.1_3.9__g8175d3d.ari     | 20) perftools-base/7.1.1                           |
| 10) ugni/6.0.14.0-7.0.1.1_7.10__ge78e5b0.ari | 21) PrgEnv-pgi/6.0.5                               |
| 11) pmi/5.0.14                               |  |

# Managing programming environments

## Switching back to the Cray programming environment

```
$ module switch PrgEnv-pgi/6.0.5 PrgEnv-cray
$ module list
Currently Loaded Modulefiles:
 1) modules/3.2.11.3
 2) slurm/20.02.2-1
 3) xalt/2.7.24
 4) cce/9.0.2
 5) craype-haswell
 6) craype-network-aries
 7) craype/2.6.1
 8) cray-mpich/7.7.10
 9) cray-libsci/19.06.1
10) udreg/2.3.2-7.0.1.1_3.9__g8175d3d.ari
11) ugni/6.0.14.0-7.0.1.1_7.10__ge78e5b0.ari
12) pmi/5.0.14
13) dmapp/7.1.1-7.0.1.1_4.8__g38cf134.ari
14) gni-headers/5.0.12.0-7.0.1.1_6.7__g3b1768f.ari
15) xpmem/2.2.19-7.0.1.1_3.7__gdcf436c.ari
16) job/2.2.4-7.0.1.1_3.8__g36b56f4.ari
17) dvs/2.12_2.2.151-7.0.1.1_5.6__g7eb5e703
18) alps/6.6.56-7.0.1.1_4.10__g2e60a7e4.ari
19) rca/2.2.20-7.0.1.1_4.9__g8e3fb5b.ari
20) atp/2.1.3
21) perftools-base/7.1.1
22) PrgEnv-cray/6.0.5
```

# Managing programming environments

## Loading and unloading modules

```
$ module load cray-hdf5
$ which h5dump
/opt/cray/pe/hdf5/1.10.5.1/bin/h5dump
$ module unload cray-hdf5
$ which h5dump
which: no h5dump in (/opt/cray/pe/perftools/7.1.1/bin:/opt/cray/pe/papi/5.7.0.2/bin:/opt/cray/rca
➤ /2.2.20-7.0.1.1_4.9__g8e3fb5b.ari/bin:/opt/cray/alps/6.6.56-7.0.1.1_4.10__g2e60a7e4.ari/sbin:/opt
➤ /cray/alps/default/bin:/opt/cray/job/2.2.4-7.0.1.1_3.8__g36b56f4.ari/bin:/opt/cray/pe/mpt/7.7.10/
➤ gni/bin:/opt/cray/pe/craype/2.6.1/bin:/opt/cray/pe/cce/9.0.2/cce-clang/x86_64/bin:/opt/cray/pe/
➤ cce/9.0.2/binutils/x86_64/x86_64-pc-linux-gnu/bin:/opt/cray/pe/cce/9.0.2/binutils/cross/x86_64-
➤ aarch64/aarch64-linux-gnu/./bin:/opt/cray/pe/cce/9.0.2/utils/x86_64/bin:/usr/karakasv/local/
➤ bin:/usr/karakasv/local/sbin:/apps/daint/UES/xalt/xalt2/software/xalt/2.7.24/sbin:/apps/daint/
➤ UES/xalt/xalt2/software/xalt/2.7.24/bin:/opt/cray/pe/modules/3.2.11.3/bin:/apps/daint/system/bin
➤ :/apps/common/system/bin:/usr/local/bin:/usr/bin:/bin:/usr/lib/mit/bin:/usr/lib/mit/sbin:/opt/
➤ cray/pe/bin)
$ h5dump
If 'h5dump' is not a typo you can use command-not-found to lookup the package that contains it, like
➤ this:
cnf h5dump
```

# Managing programming environments

## Checking available modules

```
$ module avail gcc
```

```
----- /opt/modulefiles -----  
gcc/6.1.0          gcc/7.3.0          gcc/8.1.0          gcc/8.2.0          gcc/8.3.0(default)
```

# Managing programming environments

## Loading the CSCS software stack

```
$ module load daint-gpu  
$ module avail TensorFlow
```

```
----- /apps/daint/UES/jenkins/7.0.UP01/gpu/easybuild/modules/all -----  
TensorFlow/1.14.0-CrayGNU-19.10-cuda-10.1.168-python3  
TensorFlow/2.0.0-CrayGNU-19.10-cuda-10.1.168  
TensorFlow/2.2.0-CrayGNU-19.10-cuda-10.1.168
```

# Managing programming environments

Get information about a module

```
$ module show gcc
```

```
-----  
/opt/modulefiles/gcc/8.3.0:
```

```
conflict      gcc  
conflict      gcc-cross-aarch64  
prepend-path  PATH /opt/gcc/8.3.0/bin  
prepend-path  MANPATH /opt/gcc/8.3.0/snos/share/man  
prepend-path  INFOPATH /opt/gcc/8.3.0/snos/share/info  
prepend-path  LD_LIBRARY_PATH /opt/gcc/8.3.0/snos/lib64  
setenv        GCC_PATH /opt/gcc/8.3.0  
setenv        GCC_VERSION 8.3.0  
setenv        GNU_VERSION 8.3.0  
-----
```

# Managing programming environments

Get help for a module

```
$ module help TensorFlow
```

```
----- Module Specific Help for 'TensorFlow/2.2.0-CrayGNU-19.10-cuda-10.1.168'
      └-----
```

Description

=====

An open-source software library for Machine Intelligence.

More information

=====

- Homepage: <https://www.tensorflow.org/>

Included extensions

=====

grpcio-1.24.0, tensorboard-2.2.1, tensorflow-estimator-2.2.0

# Running on Piz Daint

The job scheduler

Piz Daint uses native SLURM for running jobs on the compute nodes. There are three ways of submitting a job:

1. Interactively from the login nodes using the `srun` command.
2. By submitting a job script using the `sbatch` command.
3. By pre-allocating resources using the `salloc` command.



# Running on Piz Daint

Using the `srn` command

Necessary and useful options:

- `-C gpu`: requests allocation on the hybrid (GPU) nodes (required)
- `--reservation=summer_school`
  - 210 nodes valid until July 24 @ 20:30.
- `-N 2`: number of compute nodes (default is 1)
- `-n 2`: number of MPI tasks (default is 1)
- `-t 5`: maximum duration of the job (default is 30min)
  - Allows to get an allocation quicker
  - Job will be killed if time limit is reached
  - Maximum time slot for a job is 24h

More on <https://user.cscs.ch/access/running>

# Running on Piz Daint

Using the `srun` command

```
$ srun --reservation=summer_school -Cgpu -t1 -N2 hostname  
srun: job 24097869 queued and waiting for resources  
srun: job 24097869 has been allocated resources  
nid01994  
nid01995
```

# Running on Piz Daint

## Using the sbatch command

```
$ cat > job.sh
#!/bin/bash
#SBATCH -J 'my_first_job'
#SBATCH -C gpu
#SBATCH -N 2
#SBATCH --reservation=summer_school
#SBATCH -o myjob.out
#SBATCH -e myjob.err
```

```
echo "My job id is $SLURM_JOB_ID"
srun hostname
^D
```

```
$ sbatch job.sh
Submitted batch job 24101970
```

```
$ squeue -j 24101970
```

JOBID	USER	ACCOUNT	NAME	ST	REASON	START_TIME	TIME	TIME_LEFT	NODES	CPUS
24101970	karakasv	csstaff	my_first_job	PD	Priority	N/A	0:00	1:00	2	2

```
$ squeue -j 24101970
```

JOBID	USER	ACCOUNT	NAME	ST	REASON	START_TIME	TIME	TIME_LEFT	NODES	CPUS
24101970	karakasv	csstaff	my_first_job	R	None	18:14:26	0:21	0:39	2	48

```
$ cat myjob.err
$ cat myjob.out
My job id is 24101970
nid01985
nid01986
```

# Running on Piz Daint

Using the salloc command

```
$ salloc -Cgpu -t1 -N2
salloc: Pending job allocation 24102049
salloc: job 24102049 queued and waiting for resources
salloc: job 24102049 has been allocated resources
salloc: Granted job allocation 24102049
$ srun -N2 hostname
nid01986
nid01985
$ srun -N1 hostname
nid01985
$ srun -N4 hostname
srun: error: Only allocated 2 nodes asked for 4
$ exit
salloc: Relinquishing job allocation 24102049
salloc: Job allocation 24102049 has been revoked.
```

# Running on Piz Daint

## Other useful commands

- `squeue [OPTIONS]`: Check the status of the system job queue
  - Useful options: `-u [USERNAME]`, `-j [JOBID]`
- `scancel [JOBID]`: Cancel a job
- `scontrol`: Detailed information about partitions, reservations, computing nodes etc.

# Running on Piz Daint

## Other useful commands

```
$ scontrol show reservation summer_school
ReservationName=summer_school StartTime=Mon 08:30 EndTime=Mon 20:30 Duration=12:00:00
Nodes=nid0[3860-3887,3892-4073] NodeCnt=210 CoreCnt=2520 Features=(null) PartitionName=(null) Flags=
    ↪WEEKDAY,SPEC_NODES
TRES=cpu=5040
Users=(null) Accounts=class02,class03,class04,csstaff Licenses=(null) State=INACTIVE BurstBuffer=(
    ↪null) Watts=n/a
MaxStartDelay=(null)

[18:22:47] karakasv@daint106 [~]
$ scontrol show partition normal
PartitionName=normal
AllowGroups=ALL AllowAccounts=ALL AllowQos=ALL
AllocNodes=ALL Default=YES QoS=N/A
DefaultTime=00:30:00 DisableRootJobs=NO ExclusiveUser=NO GraceTime=0 Hidden=NO
MaxNodes=2400 MaxTime=1-00:00:00 MinNodes=1 LLN=NO MaxCPUsPerNode=UNLIMITED
Nodes=nid0
    ↪[0004-0007,0012-0024,0026-0062,0064-0067,0072-0126,0128-0190,0192-0195,0200-0254,0260-0318,0320-0323,0
    ↪
PriorityJobFactor=10 PriorityTier=20 RootOnly=NO ReqResv=NO OverSubscribe=EXCLUSIVE
OverTimeLimit=NONE PreemptMode=OFF
State=UP TotalCPUs=246720 TotalNodes=7212 SelectTypeParameters=NONE
JobDefaults=(null)
DefMemPerNode=UNLIMITED MaxMemPerNode=UNLIMITED
```

# Editing files

- `vim` or `gvim` (X version)
- `emacs -nw` or just `emacs` (X version)
- `gedit` (X only)

# Moving data to/from CSCS

- **scp**: Remote copy over SSH
  - Getting a file: `scp classXXX@ela.cscs.ch:remotefile localfile`
  - Getting a directory: `scp -r classXXX@ela.cscs.ch:remotedir localdir`
  - Sending a file: `scp localfile classXXX@ela.cscs.ch:remotefile`
  - Sending a directory: `scp localdir classXXX@ela.cscs.ch:remotedir`
  - NOTE: The `$SCRATCH` filesystem is not mounted in the frontend nodes
- **rsync**: Synchronize files remotely over SSH
  - `rsync -avz classXXX@ela.cscs.ch:remotedir/ localdir/`
  - `rsync -avz localdir/ classXXX@ela.cscs.ch:remotedir/`
  - Pay attention to the slashes! *rsync behaves differently with or without slashes.*



# Summer school repository

All the material of the course is placed inside the following Github repo:

- <https://github.com/eth-cscs/SummerSchool2020>
- For instructions on how to clone and pull from the repository, check its front page.

Organization of the repository

- `miniapp/`: The different versions of the mini-app that you will work throughout the summer school + slides.
- `topics/`: The practical exercises of the different topics that will be covered during the the summer school + slides.
- `scripts/`: Useful scripts for the exercises and the mini-app.

Solutions:

- The solutions of the exercises and the mini-app will appear at the end of the summer school in subfolders named `solutions/` under each respective topic.